

ISA vs Microarchitecture



First LC-3 Program: Use of Conditional Branches for Looping

Readings

- *Introduction to Computing Systems: From Bits and Gates to C and Beyond* by Patt & Patel
 - Chapter 5, 6
 - Appendix A – LC3 ISA

An Algorithm for Adding Integers

- We want to write a program that adds 12 integers
 - They are stored in addresses 0x3100 to 0x310B
 - Let us take a look at the flowchart of the algorithm



R1: initial address of integers

R3: final result of addition

R2: number of integers
left to be added

Check if R2 becomes 0
(done with all integers?)

Load integer in R4

Accumulate integer value in R3


Increment address R1

Decrement R2

A Program for Adding Integers in LC-3

- We use conditional branch instructions to create a loop

Address	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0	
x3000	LEA	1	1	0	0	0	1	0x00FF	1	1	1	1	1	1	1	1	R1 = PC [†] + 0x00FF = 3100 // load address
x3001	AND	1	0	1	0	1	1	0	1	1	1	0	0	0	0	0	R3 = 0 // reset register
x3002	AND	1	0	1	0	1	0	0	1	0	1	0	0	0	0	0	R2 = 0 // reset register
x3003	ADD	0	0	1	0	1	0	0	1	0	1	0	1	1	0	0	R2 = R2 + 12 // initialize counter
x3004	BR	0	0	0	0	z	0	5	0	0	0	0	0	0	1	0	BRz (PC [†] + 5) = BRz 0x300A // check condition
x3005	LDR	1	1	0	1	0	0	0	0	1	0	0	0	0	0	0	R4 = M[R1 + 0] // load value
x3006	ADD	0	0	1	0	1	1	0	1	1	0	0	0	1	0	0	R3 = R3 + R4 // accumulate
x3007	ADD	0	0	1	0	0	1	0	0	1	1	0	0	0	0	1	R1 = R1 + 1 // increment address
x3008	ADD	0	0	1	0	1	0	0	1	0	1	1	1	1	1	1	R2 = R2 - 1 // decrement counter
x3009	BR	0	0	0	n	z	p	-6	1	1	1	1	1	0	1	0	BRnzp (PC [†] - 6) = BRnzp 0x3004 // jump

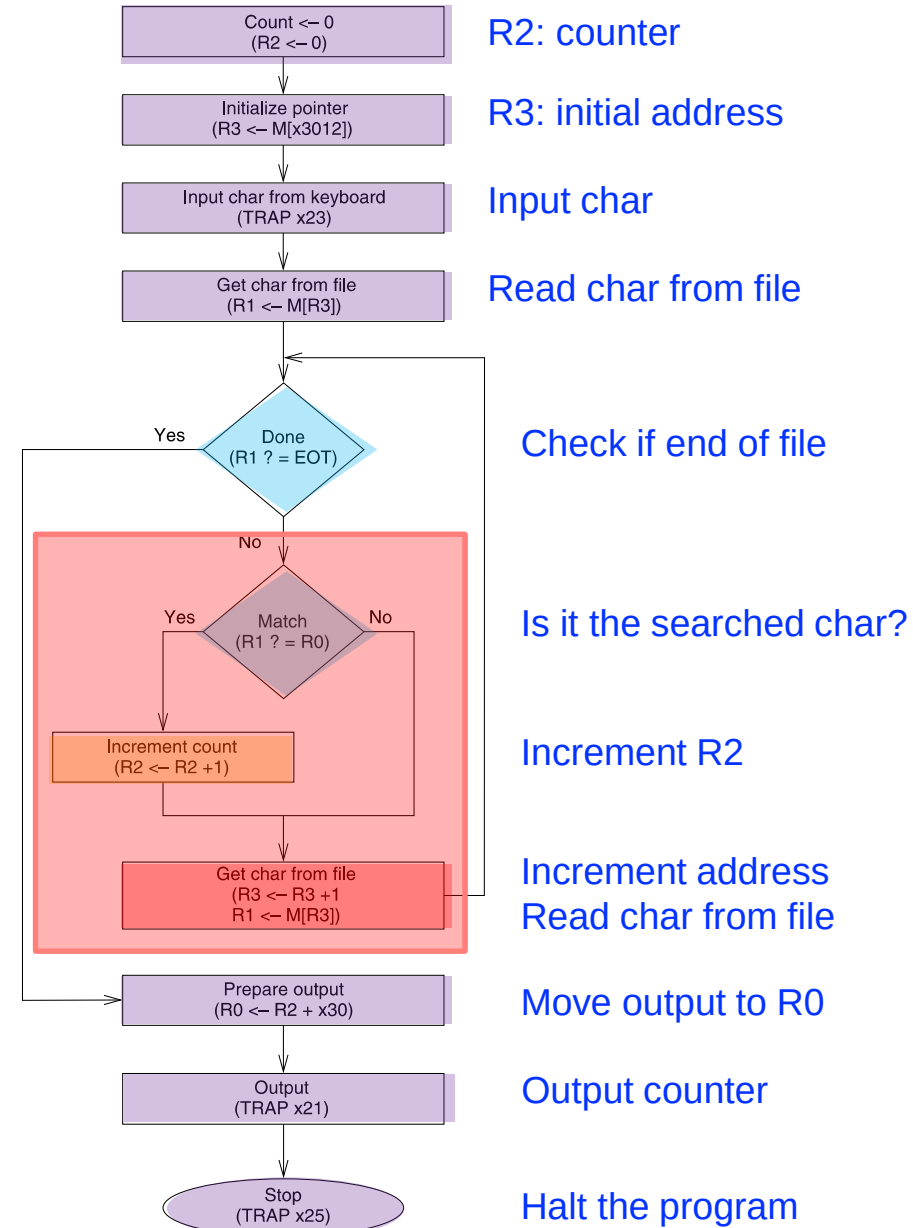


Bit 5 to differentiate both ADD instructions

[†] This is the incremented PC

Counting Occurrences of a Character

- We want to write a program that counts the occurrences of a character in a file
 - Character from the keyboard (TRAP instr.)
 - The file finishes with the character EOT (End Of Text)
 - That is called a sentinel
 - In this example, EOT = 4
 - Result to the monitor (TRAP instr.)



Counting Occurrences of a Char in LC-3

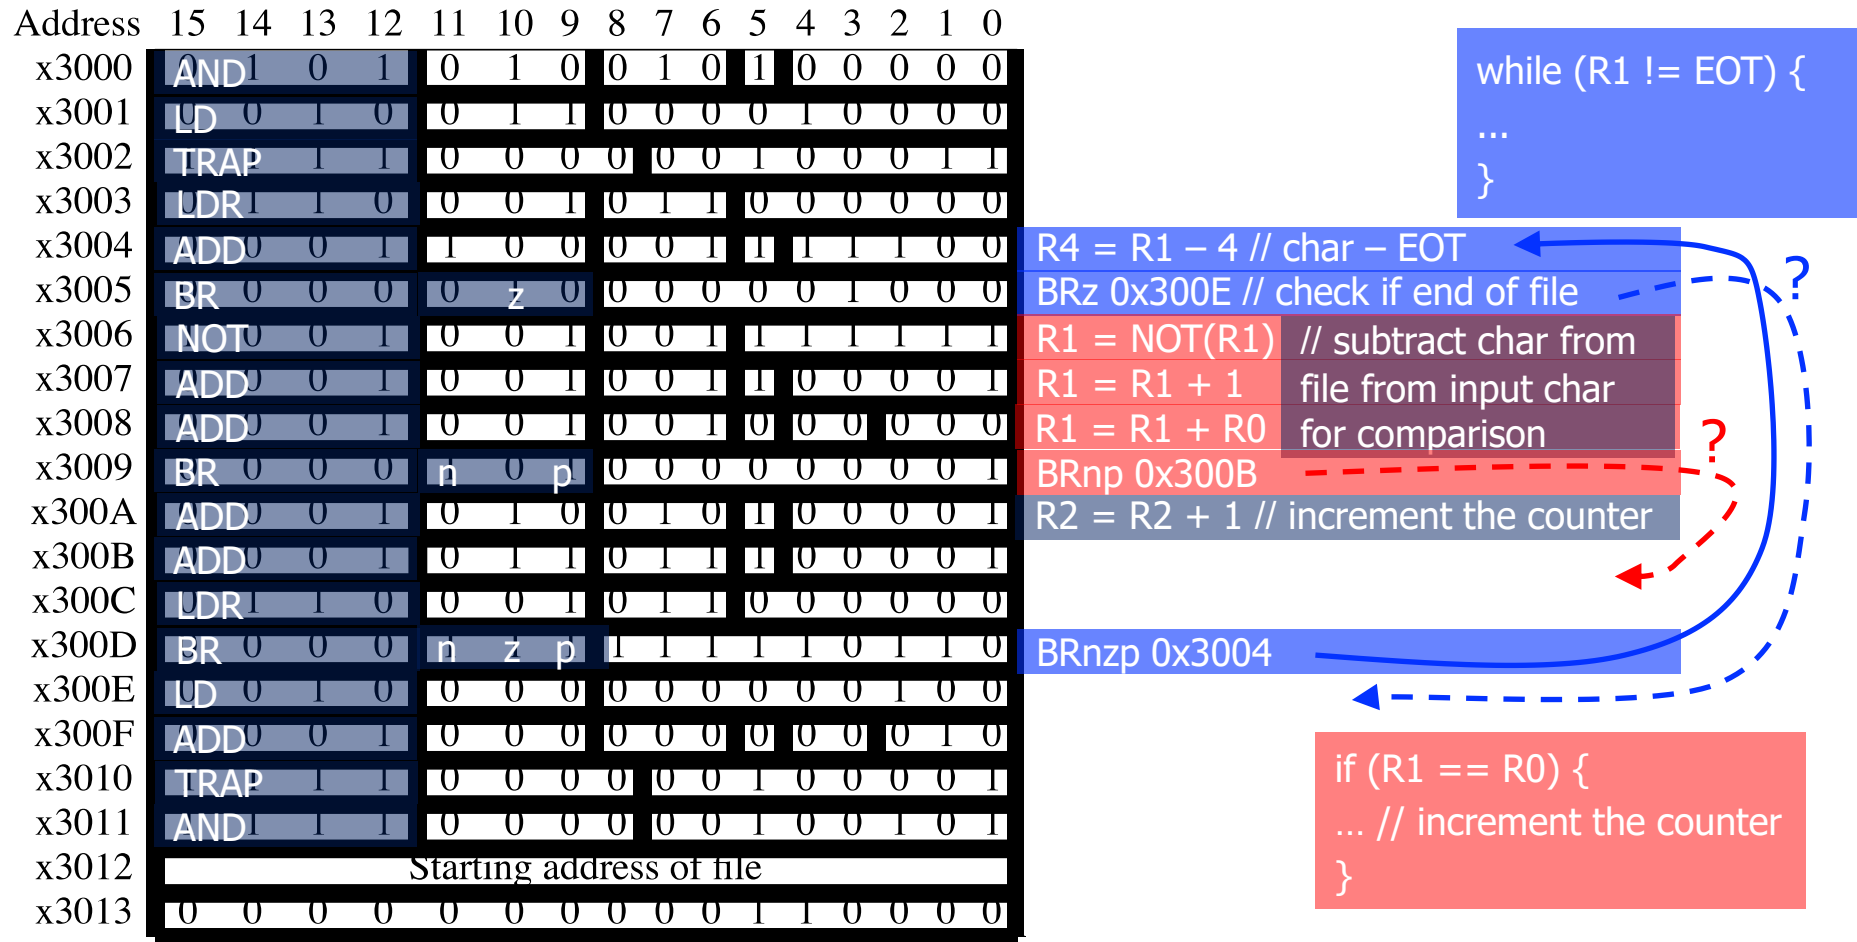
- We use **conditional branch instructions** to create a **loops** and **if statements**

Address	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0	
x3000	AND	1	0	1	0	1	0	0	1	0	1	0	0	0	0	0	R2 = 0 // initialize counter
x3001	LD	0	1	0	0	1	1	0	0	0	0	1	0	0	0	0	R3 = M[0x3012] // initial address
x3002	TRAP	1	1	0	0	0	0	0	0	1	0	0	0	0	1	1	TRAP 0x23 // input char to R0
x3003	LDR	1	1	0	0	0	1	0	1	1	0	0	0	0	0	0	R1 = M[R3] // char from file
x3004	ADD	0	0	1	1	0	0	0	0	1	1	1	1	1	0	0	R4 = R1 - 4 // char - EOT
x3005	BR	0	0	0	0	z	0	0	0	0	0	0	1	0	0	0	BRz 0x300E // check if end of file
x3006	NOT	0	0	1	0	0	1	0	0	1	1	1	1	1	1	1	R1 = NOT(R1) // subtract char from
x3007	ADD	0	0	1	0	0	1	0	0	1	1	0	0	0	0	1	file from input char
x3008	ADD	0	0	1	0	0	1	0	0	1	0	0	0	0	0	0	R1 = R1 + R0 for comparison
x3009	BR	0	0	0	n	0	p	0	0	0	0	0	0	0	0	1	BRnp 0x300B
x300A	ADD	0	0	1	0	1	0	0	1	0	1	0	0	0	0	1	R2 = R2 + 1 // increment the counter
x300B	ADD	0	0	1	0	1	1	0	1	1	1	0	0	0	0	1	R3 = R3 + 1 // increment address
x300C	LDR	1	1	0	0	0	1	0	1	1	0	0	0	0	0	0	R1 = M[R3] // char from file
x300D	BR	0	0	0	n	z	p	1	1	1	1	1	0	1	1	0	BRnzp 0x3004
x300E	LD	0	1	0	0	0	0	0	0	0	0	0	0	1	0	0	R0 = M[0x3013] // output counter to
x300F	ADD	0	0	1	0	0	0	0	0	0	0	0	0	0	1	0	monitor with T
x3010	TRAP	1	1	0	0	0	0	0	0	1	0	0	0	0	0	1	TRAP 0x21
x3011	AND	1	1	1	0	0	0	0	0	0	1	0	0	1	0	1	TRAP 0x25
x3012	Starting address of file																
x3013	ASCII TEMPLATE																

// output counter to
monitor with T
RAP

Programming Constructs in LC-3

- Let us do some reverse engineering to identify **conditional constructs** and **iterative constructs**



[illegible]

ISA vs. Microarchitecture

- ISA
 - Agreed upon interface between software and hardware
 - SW/compiler assumes, HW promises
 - What the software writer needs to know to write and debug system/user programs
- Microarchitecture
 - Specific implementation of an ISA
 - Not visible to the software
- “Computer Architecture” = ISA + microarchitecture

Problem
Algorithm
Program
ISA
Microarchitecture
Circuits
Electrons

ISA vs. Microarchitecture

- What is part of ISA vs. Uarch?
 - Gas pedal: interface for “acceleration”
 - Internals of the engine: implement “acceleration”
- Implementation (uarch) can be various as long as it satisfies the specification (ISA)
 - Add instruction vs. Adder implementation
 - Bit serial, ripple carry, carry lookahead adders are all part of microarchitecture
 - x86 ISA has many implementations: 286, 386, 486, Pentium, Pentium Pro, Pentium 4, Core, ...
- Microarchitecture usually changes faster than ISA
 - Few ISAs (x86, ARM, SPARC, MIPS, Alpha) but many uarchs
 - *Why?*

ISA

- Instructions
 - Opcodes, Addressing Modes, Data Types
 - Instruction Types and Formats
 - Registers, Condition Codes
- Memory
 - Address space, Addressability, Alignment
 - Virtual memory management
- Interrupt/Exception Handling
- Access Control, Priority/Privilege
- I/O: memory-mapped vs. special instruction
- Task/thread Management
- Multi-threading support, Multiprocessor support

Microarchitecture

- Implementation of the ISA under specific **design constraints and goals**
- Anything done in hardware without exposure to software
 - Pipelining
 - In-order versus out-of-order instruction execution
 - Memory access scheduling policy
 - Speculative execution
 - Superscalar processing
 - Clock gating
 - Caching - Levels, size, associativity, replacement policy
 - Prefetching

Property of ISA vs. Uarch?

- ADD instruction's opcode?
- Number of general purpose registers?
- Number of ports to the register file?
- Number of cycles to execute the MUL instruction?
- Whether or not the machine employs pipelined instruction execution?
- Remember
 - Microarchitecture: Implementation of the ISA under specific design constraints and goals

Design Point

- A set of design considerations and their importance
 - **leads to tradeoffs** in both ISA and uarch
- Considerations
 - Cost
 - Performance
 - Maximum power consumption
 - Energy consumption (battery life)
 - Availability
 - Reliability and Correctness
 - Time to Market
- Design point determined by the “Problem” space (application space), the intended users/*market*

Problem
Algorithm
Program
ISA
Microarchitecture
Circuits
Electrons


Application Space

- Scientific applications
 - Predicting weather
 - Predicting earthquake
- Transaction-based applications
 - ATM transfers
 - E-commerce business
- Network applications
 - High-speed routing of Internet packets
- Real-time applications
 - Meeting critical deadline
- IoT applications
 - Controlling small sensor devices
- Media applications
 - Decoding video and audio files
- Cost
- Performance
- Maximum power consumption
- Energy consumption (battery life)
- Availability
- Reliability and Correctness
- Time to Market

Tradeoffs: Soul of Computer Architecture

- ISA-level tradeoffs
- Microarchitecture-level tradeoffs
- System and Task-level tradeoffs
 - How to divide the labor between hardware and software
- *Computer architecture is the science and art of making the appropriate trade-offs to meet a design point*
 - *Why art?*

Why Is It (Somewhat) Art?



We do not (fully) know the future (applications, users, market)

Many Different ISAs Over Decades

- x86
- PDP-x: Programmed Data Processor (PDP-11)
- VAX
- IBM 360
- CDC 6600
- SIMD ISAs: CRAY-1, Connection Machine
- VLIW ISAs: Multiflow, Cydrome, IA-64 (EPIC)
- RISC ISAs: Alpha, MIPS, SPARC, ARM
- What are the fundamental differences?
 - E.g., how instructions are specified and what they do
 - E.g., how complex are the instructions

What Are the Elements of An ISA?

- Instructions

- Opcode
- Operand specifiers (addressing modes)
 - How to obtain the operand?

- Data types

- Definition: Representation of information for which there are instructions that operate on the representation
- Integer, floating point, character, binary, decimal, BCD
- Doubly linked list, queue, string, bit vector, stack

Data Type Tradeoffs

- What is the benefit of having more or high-level data types in the ISA?
- What is the disadvantage?
- Think compiler/programmer vs. microarchitect
- Concept of semantic gap
 - Data types coupled tightly to the semantic level, or complexity of instructions
- Example: Early RISC architectures vs. Intel 432
 - Early RISC: Only integer data type
 - Intel 432: Object data type

What Are the Elements of An ISA?


- Memory organization

- Address space: How many uniquely addressable locations in memory
- Addressability: How much data does each uniquely identifiable location store
 - Byte addressable: most ISAs, characters are 8 bits
 - 32-bit addressable: First Alpha
- Support for virtual memory

What Are the Elements of An ISA?


- Registers
 - How many
 - Size of each register
- Why is having registers a good idea?
 - Because programs exhibit a characteristic called **data locality**
 - **A recently produced/accessed value is likely to be used more than once (temporal locality)**
 - Storing that value in a register eliminates the need to go to memory each time that value is needed

Programmer Visible (Architectural) State




Memory

Array of storage locations indexed by an address



Registers

- given special names in the ISA
(as opposed to addresses)
- general vs. special purpose



memory address of the current instruction

Instructions specify how to transform the values of programmer visible state

Sequential FSM

Aside: Programmer Invisible State

- Microarchitectural state
- Programmer cannot access this directly
- E.g. cache state

Instruction Classes

- Operate instructions
 - Process data: arithmetic and logical operations
 - Fetch operands, compute result, store result
 - Implicit sequential control flow
- Data movement instructions
 - Move data between memory, registers, I/O devices
 - Implicit sequential control flow
- Control flow instructions
 - Change the sequence of instructions that are executed

What Are the Elements of An ISA?


- Load/store vs. memory/memory architectures
 - Load/store architecture: operate instructions operate only on registers
 - E.g., MIPS, ARM and many RISC ISAs
 - Memory/memory architecture: operate instructions can operate on memory locations
 - E.g., x86, VAX and many CISC ISAs

What Are the Benefits of Different Addressing Modes?

- Another example of programmer vs. microarchitect tradeoff
- Advantage of more addressing modes:
 - Enables **better mapping of high-level constructs to the machine**: some accesses are better expressed with a different mode → reduced number of instructions and code size
 - Think array accesses (autoincrement mode)
 - Think indirection (pointer chasing)
 - Sparse matrix accesses
- Disadvantage:
 - **More work** for the **compiler**
 - **More work** for the **microarchitect**

What Are the Elements of An ISA?

- How to interface with I/O devices
 - Memory mapped I/O
 - A region of memory is mapped to I/O devices
 - I/O operations are loads and stores to those locations
 - Special I/O instructions
 - IN and OUT instructions in x86 deal with ports of the chip
 - Tradeoffs?
 - Which one is more general purpose?



What Are the Elements of An ISA?

- Privilege modes
 - User vs supervisor
 - Who can execute what instructions?
- Exception and interrupt handling
 - What procedure is followed when something goes wrong with an instruction?
 - What procedure is followed when an external device requests the processor?
 - Vectored vs. non-vectored interrupts (early MIPS)
- Virtual memory
 - Each program has the illusion of the entire memory space, which is greater than physical memory
- Access protection
- We will talk about these later

Complex vs. Simple Instructions

- Complex instruction: An instruction does a lot of work, e.g. many operations
 - Insert in a doubly linked list
 - INDEX instruction (array access with bounds checking)
 - String copy
- Simple instruction: An instruction does small amount of work, it is a primitive using which complex operations can be built
 - Add
 - XOR
 - Multiply

Complex vs. Simple Instructions

- Advantages of Complex instructions
 - + Denser encoding → smaller code size → better memory utilization, better cache hit rate (better packing of instructions)
 - + Simpler compiler: no need to optimize small instructions as much
- Disadvantages of Complex Instructions
 - Larger chunks of work → compiler has less opportunity to optimize (limited in fine-grained optimizations it can do)
 - More complex hardware → translation from a high level to control signals and optimization needs to be done by hardware

ISA-level Tradeoffs: Semantic Gap

- Where to place the ISA? Semantic gap
 - Closer to high-level language (HLL) → Small semantic gap, complex instructions
 - Closer to hardware control signals? → Large semantic gap, simple instructions
- RISC vs. CISC machines
 - RISC: Reduced instruction set computer
 - CISC: Complex instruction set computer
 - QUICKSORT, FP instructions?
 - VAX INDEX instruction (array access with bounds checking)

ISA-level Tradeoffs: Semantic Gap

- Some tradeoffs
- Simple compiler, complex hardware vs. complex compiler, simple hardware
 - Caveat: Translation (indirection) can change the tradeoff!
- Burden of backward compatibility
- Performance?
 - Optimization opportunity: Example of VAX INDEX instruction: who (compiler vs. hardware) puts more effort into optimization?
 - Instruction size, code size

Small versus Large Semantic Gap

- CISC vs. RISC
 - Complex instruction set computer → complex instructions
 - Initially motivated by “not good enough” code generation
 - Reduced instruction set computer → simple instructions
 - John Cocke, mid 1970s, IBM 801
 - Goal: enable better compiler control and optimization
- RISC motivated by
 - Simplifying the hardware → lower cost, higher frequency
 - Enabling the compiler to optimize the code better
 - Find fine-grained parallelism

How High or Low Can You Go?

- Very large semantic gap
 - Each instruction specifies the complete set of control signals in the machine
 - Compiler generates control signals
 - Open microcode (John Cocke, circa 1970s)
 - Gave way to optimizing compilers
- Very small semantic gap
 - ISA is (almost) the same as high-level language
 - Java machines, LISP machines, object-oriented machines

A Note on ISA Evolution

- ISAs have evolved to reflect/satisfy the concerns of the day
- Examples:
 - Limited compiler optimization technology
 - Limited memory bandwidth
 - Need for specialization in important applications
- Use of translation (in HW and SW) enabled underlying implementations to be similar, regardless of the ISA

Effect of Translation

- One can translate from one ISA to another *ISA* to change the semantic gap tradeoffs
 - ISA (virtual ISA) → Implementation ISA
- Examples
 - Intel's and AMD's x86 implementations translate x86 instructions into programmer-invisible microoperations (simple instructions) in hardware
 - Transmeta's x86 implementations translated x86 instructions into "secret" VLIW instructions in software (code morphing software)

Hardware-Based Translation




Figure 2. Conventional superscalar out-of-order CPUs use hardware to create and dispatch micro-ops that can execute in parallel.

Klaiber, [“The Technology Behind Crusoe Processors,”](#) Transmeta White Paper 2000.

Software-Based Translation




Figure 5. The Code Morphing software mediates between x86 software and the Crusoe processor.

Klaiber, "[The Technology Behind Crusoe Processors](#)," Transmeta White Paper 2000.

ISA-level Tradeoffs: Instruction Length


- **Fixed length:** Length of all instructions the same
 - + Easier to decode single instruction in hardware
 - + Easier to decode multiple instructions concurrently
 - Wasted bits in instructions
 - Harder-to-extend ISA (how to add new instructions?)
- **Variable length:** Length of instructions different (determined by opcode and sub-opcode)
 - + Compact encoding
 - Intel 432: Huffman encoding (sort of). 6 to 321 bit instructions. **How?**
 - More logic to decode a single instruction
 - Harder to decode multiple instructions concurrently
- **Tradeoffs**
 - Code size (memory space, bandwidth) vs. hardware complexity
 - ISA extensibility and expressiveness vs. hardware complexity
 - Energy? Smaller code vs. ease of decode

ISA-level Tradeoffs: Uniform Decode


- **Uniform decode:** Same bits in each instruction correspond to the same meaning
 - Opcode is always in the same location
 - Ditto operand specifiers, immediate values, ...
 - Many “RISC” ISAs: Alpha, MIPS, SPARC
 - + Easier decode, simpler hardware
 - Wastes space
- **Non-uniform decode**
 - E.g., opcode can be the 1st-7th byte in x86
 - + More compact and powerful instruction format
 - More complex decode logic

x86 vs. Alpha Instruction Formats

- x86:




- Alpha:



MIPS Instruction Format


- R-type, 3 register operands



- I-type, 2 register operands and 16-bit immediate operand



- J-type, 26-bit immediate operand



- Simple Decoding

- 4 bytes per instruction, regardless of format
- must be 4-byte aligned (2 lsb of PC must be '00')
- format and fields easy to extract in hardware

A Note on RISC vs. CISC

- Usually, ...
- RISC
 - Simple instructions
 - Fixed length
 - Uniform decode
 - Few addressing modes
- CISC
 - Complex instructions
 - Variable length
 - Non-uniform decode
 - Many addressing modes

ISA-level Tradeoffs: Number of Registers

- Affects:
 - Number of bits used for encoding register address
 - Number of values kept in fast storage (register file)
 - (uarch) Size, access time, power consumption of register file
- Large number of registers:
 - + Enables better register allocation (and optimizations) by compiler → fewer saves/restores
 - Larger instruction size

ISA-level Tradeoffs: Addressing Modes

- Addressing mode specifies how to obtain an operand of an instruction
 - Register
 - Immediate
 - Memory (displacement, register indirect, indexed, absolute, memory indirect, autoincrement, autodecrement, ...)
- More modes:
 - + help better support programming constructs (arrays, pointer-based accesses)
 - make it harder for the architect to design
 - too many choices for the compiler?
 - Many ways to do the same thing complicates compiler design

Back to Programmer vs. (Micro)architect

- Many ISA features designed to aid programmers
- But, complicate the hardware designer's job
- Virtual memory
 - vs. overlay programming
 - Should the programmer be concerned about the size of code blocks fitting physical memory?
- Addressing modes
- Unaligned memory access
 - Compiler/programmer needs to align data