

INHA UNIVERSITY TASHKENT

FALL SEMESTER 2024

SOC 3010

OPERATING SYSTEM

CREDITS/HOURS PER WEEK : 3/3

COURSE TYPE : TECHNICAL CORE SEQUENCE

INSTRUCTOR

DR. A. R. NASEER

HEAD & PROFESSOR OF COMPUTER SCIENCE & ENGG

SOC 3010

OPERATING SYSTEM

WEEK 3 LECTURE1

UNIX KERNEL

OVERVIEW

Overview of Unix Kernels

- ★ Unix kernels provide an execution environment in which applications may run.
- ★ Therefore, the kernel must implement a set of services and corresponding interfaces.
- ★ Applications use those interfaces and do not usually interact directly with hardware resources.

Process/Kernel Model

- ★ A OS can operate either in User Mode or in Kernel Mode
- ★ All standard Unix kernels make use of only Kernel Mode and User Mode.
- ★ When a program is executed in User Mode, it cannot directly access the kernel data structures or the kernel programs.
- ★ When an application executes in Kernel Mode, however, these restrictions no longer apply.
- ★ Each CPU model provides special instructions to switch from User Mode to Kernel Mode and vice versa.
- ★ A program executes most of the time in User Mode and switches to Kernel Mode only when requesting a service provided by the kernel.
- ★ When the kernel has satisfied the program's request, it puts the program back in User Mode.

Process/Kernel Model

- ★ Processes are dynamic entities that usually have a limited life span within the system.
- ★ The task of creating, eliminating, and synchronizing the existing processes is delegated to a group of routines in the kernel.
- ★ The kernel itself is not a process but a process manager.
- ★ The process/kernel model assumes that processes that require a kernel service make use of specific programming constructs called *system calls*.
- ★ Each system call sets up the group of parameters that identifies the process request and then executes the hardware-dependent CPU instruction to switch from User Mode to Kernel Mode.


Process/Kernel Model

Besides user processes, Unix systems include a few privileged processes called *kernel threads* with the following characteristics:

- They run in Kernel Mode in the kernel address space.
- They do not interact with users, and thus do not require terminal devices.
- They are usually created during system startup and remain alive until the system is shut down.

Process/Kernel Model

Example: Transitions between User and Kernel Mode



- * Process 1 in User Mode issues a system call, after which the process switches to Kernel Mode and the system call is serviced.
- * Process 1 then resumes execution in User Mode until a timer interrupt occurs and the scheduler is activated in Kernel Mode.
- * A process switch takes place, and Process 2 starts its execution in User Mode until a hardware device raises an interrupt.
- * As a consequence of the interrupt, Process 2 switches to Kernel Mode and services the interrupt

Process/Kernel Model

In Unix kernels, kernel routines can be activated in several ways:

- A process invokes a system call.
- The CPU executing the process signals an *exception*, which is some unusual condition such as an invalid instruction. The kernel handles the exception on behalf of the process that caused it.
- A peripheral device issues an *interrupt signal* to the CPU to notify it of an event such as a request for attention, a status change, or the completion of an I/O operation.
 - Each interrupt signal is dealt by a kernel program called an *interrupt handler*.
 - Since peripheral devices operate asynchronously with respect to the CPU, interrupts occur at unpredictable times.
- A kernel thread is executed; since it runs in Kernel Mode, the corresponding program must be considered part of the kernel, albeit encapsulated in a process.

Process Implementation

- ★ For the kernel to manage processes, each process is represented by a *process descriptor* that includes information about the current state of the process.
- ★ When the kernel stops the execution of a process, it saves the current contents of several processor registers in the process descriptor.
- ★ These include:
 - The program counter (PC) and stack pointer (SP) registers
 - The general-purpose registers
 - The floating point registers
 - The processor control registers (Processor Status Word) containing information about the CPU state
 - The memory management registers used to keep track of the RAM accessed by the process

Process Implementation

- ★ When the kernel decides to resume executing a process, it uses the proper process descriptor fields to load the CPU registers.
- ★ Since the stored value of the program counter points to the instruction following the last instruction executed, the process resumes execution from where it was stopped.
- ★ When a process is not executing on the CPU, it is waiting for some event.
- ★ Unix kernels distinguish many wait states, which are usually implemented by queues of process descriptors; each queue corresponds to the set of processes waiting for a specific event.

SOC 3010

OPERATING SYSTEM

WEEK 3 LECTURE 1

RE-ENTRANT KERNELS

Reentrant Kernels

- * All Unix kernels are *reentrant* : this means that several processes may be executing in Kernel Mode at the same time.
- * On uniprocessor systems only one process can progress, but many of them can be blocked in Kernel Mode waiting for the CPU or the completion of some I/O operation.
- * For instance, after issuing a read to a disk on behalf of some process, the kernel will let the disk controller handle it and will resume executing other processes.
- * An interrupt notifies the kernel when the device has satisfied the read, so the former process can resume the execution.
- * One way to provide reentrancy is to write functions so that they modify only local variables and do not alter global data structures. Such functions are called *reentrant functions*.
- * But a reentrant kernel is not limited just to such reentrant functions. Instead, the kernel can include nonreentrant functions and use locking mechanisms to ensure that only one process can execute a nonreentrant function at a time.

Reentrant Kernels

- ★ Every process in Kernel Mode acts on its own set of memory locations and cannot interfere with the others.
- ★ If a hardware interrupt occurs, a reentrant kernel is able to suspend the current running process even if that process is in Kernel Mode.
- ★ This capability is very important, since it improves the throughput of the device controllers that issue interrupts.
- ★ Once a device has issued an interrupt, it waits until the CPU acknowledges it.
- ★ If the kernel is able to answer quickly, the device controller will be able to perform other tasks while the CPU handles the interrupt.

Reentrant Kernels


- ❑ Kernel reentrancy and its impact on the organization of the kernel
- * A **kernel control path** denotes the sequence of instructions executed by the kernel to handle a system call, an exception, or an interrupt.
- * In the simplest case, the CPU executes a kernel control path sequentially from the first instruction to the last.
- * When one of the following events occurs, however, the CPU interleaves the kernel control paths:

EVENT 1:

- A process executing in User Mode invokes a system call and the corresponding kernel control path verifies that the request cannot be satisfied immediately; it then invokes the scheduler to select a new process to run. As a result, a process switch occurs.
- * The first kernel control path is left unfinished and the CPU resumes the execution of some other kernel control path.
- * In this case, the two control paths are executed on behalf of two different processes.

Reentrant Kernels

□ Interleaving of kernel control paths



INTERLEAVING OF KERNEL CONTROL PATH

Reentrant Kernels

- ❑ Kernel reentrancy and its impact on the organization of the kernel

EVENT 2:

- The CPU detects an exception—for example, an access to a page not present in RAM—while running a kernel control path.
- * The first control path is suspended, and the CPU starts the execution of a suitable procedure.
- * In our example, this type of procedure could allocate a new page for the process and read its contents from disk.
- * When the procedure terminates, the first control path can be resumed.
- * In this case, the two control paths are executed on behalf of the same process.

Reentrant Kernels

- ❑ Kernel reentrancy and its impact on the organization of the kernel

EVENT 3:

- A hardware interrupt occurs while the CPU is running a kernel control path with the interrupts enabled.
- * The first kernel control path is left unfinished and the CPU starts processing another kernel control path to handle the interrupt.
- * The first kernel control path resumes when the interrupt handler terminates.
- * In this case the two kernel control paths run in the execution context of the same process and the total elapsed system time is accounted to it.
- * However, the interrupt handler doesn't necessarily operate on behalf of the process

Process Address Space

- ★ Each process runs in its private address space.
- ★ A process running in User Mode refers to private stack, data, and code areas.
- ★ When running in Kernel Mode, the process addresses the kernel data and code area and makes use of another stack.
- ★ Since the kernel is reentrant, several kernel control paths—each related to a different process—may be executed in turn.
- ★ In this case, each kernel control path refers to its own private kernel stack.
- ★ While it appears to each process that it has access to a private address space, there are times when part of the address space is shared among processes.

Process Address Space

- * In some cases Address space sharing is explicitly requested by processes; in others it is done automatically by the kernel to reduce memory usage.
- * If the same program, say an editor, is needed simultaneously by several users, the program will be loaded into memory only once, and its instructions can be shared by all of the users who need it.
- * Its data, of course, must not be shared, because each user will have separate data.
- * This kind of shared address space is done automatically by the kernel to save memory.
- * Processes can also share parts of their address space as a kind of interprocess communication, using the "shared memory" technique introduced in System V and supported by Linux.
- * Finally, Linux supports the `mmap( )` system call, which allows part of a file or the memory residing on a device to be mapped into a part of a process address space.
- * Memory mapping can provide an alternative to normal reads and writes for transferring data.
- * If the same file is shared by several processes, its memory mapping is included in the address space of each of the processes that share it.

SOC 3010

OPERATING SYSTEM

WEEK 3 LECTURE 1

SYNCHRONIZATION & CRITICAL REGIONS

Synchronization and Critical Regions

- ✱ Implementing a reentrant kernel requires the use of synchronization.
- ✱ if a kernel control path is suspended while acting on a kernel data structure, no other kernel control path will be allowed to act on the same data structure unless it has been reset to a consistent state.
- ✱ Otherwise, the interaction of the two control paths could corrupt the stored information.
- ✱ For example, let's suppose that a global variable *V* contains the number of available items of some system resource.
- ✱ A first kernel control path *A* reads the variable *V* and determines that there is just one available item.
- ✱ At this point, another kernel control path *B* is activated and reads the same variable *V*, which still contains the value 1.
- ✱ Thus, *B* decrements *V* and starts using the resource item. Then *A* resumes the execution; because it has already read the value of *V*, it assumes that it can decrement *V* and take the resource item, which *B* already uses.
- ✱ As a final result, *V* contains -1, and two kernel control paths are using the same resource item with potentially disastrous effects.

Synchronization and Critical Regions

- * When the outcome of some computation depends on how two or more processes are scheduled, the code is incorrect: we say that there is a *race condition*.
- * In general, safe access to a global variable is ensured by using *atomic operations*.
- * In the previous example, data corruption would not be possible if the two control paths read and decrement V with a single, noninterruptible operation.
- * However, kernels contain many data structures that cannot be accessed with a single operation.
- * For example, it usually isn't possible to remove an element from a linked list with a single operation, because the kernel needs to access at least two pointers at once.
- * Any section of code that should be finished by each process that begins it before another process can enter it is called a *critical region*.
- * Synchronization problems occur not only among kernel control paths but also among processes sharing common data.

Synchronization and Critical Regions

❑ POSSIBLE APPROACHES PROPOSED TO RESOLVE SYNCHRONIZATION PROBLEM

- ★ **Nonpreemptive Kernels**
- ★ **Interrupt disabling**
- ★ **Semaphores**
- ★ **Spin locks**

Nonpreemptive kernels

- * In Unix, simple solution to synchronization problems is provided by making Unix kernels nonpreemptive: when a process executes in Kernel Mode, it cannot be arbitrarily suspended and substituted with another process.
- * Therefore, on a uniprocessor system all kernel data structures that are not updated by interrupts or exception handlers are safe for the kernel to access.
- * Of course, a process in Kernel Mode can voluntarily relinquish the CPU, but in this case it must ensure that all data structures are left in a consistent state.
- * Moreover, when it resumes its execution, it must recheck the value of any previously accessed data structures that could be changed.
- * Nonpreemptability is ineffective in multiprocessor systems, since two kernel control paths running on different CPUs could concurrently access the same data structure.

Interrupt disabling

- ★ Another synchronization mechanism for uniprocessor systems consists of disabling all hardware interrupts before entering a critical region and reenabling them right after leaving it.
- ★ This mechanism, while simple, is far from optimal.
- ★ If the critical region is large, interrupts can remain disabled for a relatively long time, potentially causing all hardware activities to freeze.
- ★ Moreover, on a multiprocessor system this mechanism doesn't work at all.
- ★ There is no way to ensure that no other CPU can access the same data structures updated in the protected critical region

Semaphores

- * A widely used mechanism, effective in both uniprocessor and multiprocessor systems, relies on the use of *semaphores*.
- * A semaphore is simply a counter associated with a data structure; the semaphore is checked by all kernel threads before they try to access the data structure.
- * Each semaphore may be viewed as an object composed of:
 - An integer variable
 - A list of waiting processes
 - Two atomic methods: `down( )` and `up( )`
- * The `down( )` method decrements the value of the semaphore. If the new value is less than 0, the method adds the running process to the semaphore list and then blocks (i.e., invokes the scheduler).
- * The `up( )` method increments the value of the semaphore and, if its new value is greater than or equal to 0, reactivates one or more processes in the semaphore list.

Semaphores

- ★ Each data structure to be protected has its own semaphore, which is initialized to 1.
- ★ When a kernel control path wishes to access the data structure, it executes the down() method on the proper semaphore.
- ★ If the value of the new semaphore isn't negative, access to the data structure is granted.
- ★ Otherwise, the process that is executing the kernel control path is added to the semaphore list and blocked.
- ★ When another process executes the up() method on that semaphore, one of the processes in the semaphore list is allowed to proceed.

Spin locks

- ★ In multiprocessor systems, semaphores are not always the best solution to the synchronization problems.
- ★ Some kernel data structures should be protected from being concurrently accessed by kernel control paths that run on different CPUs.
- ★ In this case, if the time required to update the data structure is short, a semaphore could be very inefficient.
- ★ To check a semaphore, the kernel must insert a process in the semaphore list and then suspend it.
- ★ Since both operations are relatively expensive, in the time it takes to complete them, the other kernel control path could have already released the semaphore.
- ★ In these cases, multiprocessor operating systems make use of *spin locks*.

Spin locks

- ★ A spin lock is very similar to a semaphore, but it has no process list: when a process finds the lock closed by another process, it "spins" around repeatedly, executing a tight instruction loop until the lock becomes open.
- ★ Of course, spin locks are useless in a uniprocessor environment.
- ★ When a kernel control path tries to access a locked data structure, it starts an endless loop.
- ★ Therefore, the kernel control path that is updating the protected data structure would not have a chance to continue the execution and release the spin lock.
- ★ The final result is that the system hangs

Avoiding deadlocks

- * Processes or kernel control paths that synchronize with other control paths may easily enter in a *deadlocked* state.
- * The simplest case of deadlock occurs when process *p1* gains access to data structure *a* and process *p2* gains access to *b*, but *p1* then waits for *b* and *p2* waits for *a*.
- * Other more complex cyclic waitings among groups of processes may also occur.
- * Of course, a deadlock condition causes a complete freeze of the affected processes or kernel control paths.
- * As far as kernel design is concerned, deadlock becomes an issue when the number of kernel semaphore types used is high.
- * In this case, it may be quite difficult to ensure that no deadlock state will ever be reached for all possible ways to interleave kernel control paths.
- * Several operating systems, including Linux, avoid this problem by introducing a very limited number of semaphore types and by requesting semaphores in an ascending order.

SOC 3010

OPERATING SYSTEM

WEEK 3 LECTURE 1

SIGNALS & INTER-PROCESS COMMUNICATION

Signals and Interprocess Communication

- * Unix signals provide a mechanism for notifying processes of system events.
- * Each event has its own signal number, which is usually referred to by a symbolic constant such as SIGTERM.
- * There are two kinds of system events:
- * *Asynchronous notifications*
 - For instance, a user can send the interrupt signal SIGTERM to a foreground process by pressing the interrupt keycode (usually, CTRL-C) at the terminal.
- * *Synchronous errors or exceptions*
 - For instance, the kernel sends the signal SIGSEGV to a process when it accesses a memory location at an illegal address.
- * The POSIX standard defines about 30 different signals, two of which are user-definable and may be used as a primitive mechanism for communication and synchronization among processes in User Mode.

Signals and Interprocess Communication

- * In general, a process may react to a signal reception in two possible ways:
 - Ignore the signal.
 - Asynchronously execute a specified procedure (the signal handler).
- * If the process does not specify one of these alternatives, the kernel performs a *default action* that depends on the signal number.
- * The five possible default actions are:
 - Terminate the process.
 - Write the execution context and the contents of the address space in a file (*core dump*) and terminate the process.
 - Ignore the signal.
 - Suspend the process.
 - Resume the process's execution, if it was stopped.

Signals and Interprocess Communication

- * Kernel signal handling is rather elaborate since the POSIX semantics allows processes to temporarily block signals.
- * Moreover, a few signals such as SIGKILL cannot be directly handled by the process and cannot be ignored.
- * AT&T's Unix System V introduced other kinds of interprocess communication among processes in User Mode, which have been adopted by many Unix kernels:
 - semaphores, message queues, and shared memory*
- * They are collectively known as *System V IPC*.
- * The kernel implements these constructs as *IPC resources*: a process acquires a resource by invoking a `semget( )`, `msgget( )`, or `shmget( )`, system call.
- * Just like files, IPC resources are persistent: they must be explicitly deallocated by the creator process, by the current owner, or by a superuser process.

Signals and Interprocess Communication

- ★ **Semaphores** are reserved for processes in User Mode.
- ★ **Message queues** allow processes to exchange messages by making use of the **msgsnd()** and **msgget()** **system calls**, which respectively insert a message into a specific message queue and extract a message from it.
- ★ **Shared memory** provides the fastest way for processes to exchange and share data.
- ★ A process starts by issuing a **shmget()** **system call** to create a new shared memory having a required size.
- ★ After obtaining the IPC resource identifier, the process invokes the **shmat()** **system call**, which returns the starting address of the new region within the process address space.
- ★ When the process wishes to detach the shared memory from its address space, it invokes the **shmdt()** **system call**.
- ★ The implementation of shared memory depends on how the kernel implements process address spaces.

SOC 3010

OPERATING SYSTEM

WEEK 3 LECTURE 1

PROCESS MANAGEMENT

Process Management

- ★ Unix makes a neat distinction between the process and the program it is executing.
- ★ To that end, the **fork()** and **exit()** system calls are used respectively to create a new process and to terminate it, while an **exec()**-like system call is invoked to load a new program.
- ★ After such a system call has been executed, the process resumes execution with a brand new address space containing the loaded program.
- ★ The process that invokes a **fork()** is the *parent* while the new process is its *child* .

Process Management

- ★ Parents and children can find each other because the data structure describing each process includes a pointer to its immediate parent and pointers to all its immediate children.
- ★ A naive implementation of the `fork( )` would require both the parent's data and the parent's code to be duplicated and assign the copies to the child. This would be quite time-consuming.
- ★ Current kernels that can rely on hardware paging units follow the **Copy-On-Write approach**, which defers page duplication until the last moment (i.e., until the parent or the child is required to write into a page).
- ★ The `exit( )` system call terminates a process. The kernel handles this system call by releasing the resources owned by the process and sending the parent process a **SIGCHLD signal, which is ignored by default.**

Zombie processes

- ❑ **How can a parent process inquire about termination of its children?**
- ★ The `wait( )` system call allows a process to wait until one of its children terminates; it returns the process ID (PID) of the terminated child.
- ★ When executing this system call, the kernel checks whether a child has already terminated.
- ★ **A special zombie process state is introduced to represent terminated processes: a process remains in that state until its parent process executes a `wait( )` system call on it.**
- ★ The system call handler extracts some data about resource usage from the process descriptor fields; the process descriptor may be released once the data has been collected.
- ★ If no child process has already terminated when the `wait( )` system call is executed, the kernel usually puts the process in a wait state until a child terminates.
- ★ Many kernels also implement a **`waitpid( )` system call, which allows a process to wait for a specific child process.** Other variants of `wait( )` system calls are also quite common.

Zombie processes

- ★ It's a good practice for the kernel to keep around information on a child process until the parent issues its `wait( )` call, but suppose the parent process terminates without issuing that call?
- ★ The information takes up valuable memory slots that could be used to serve living processes.
- ★ For example, many shells allow the user to start a command in the background and then log out.
- ★ The process that is running the command shell terminates, but its children continue their execution.
- ★ The solution lies in a special system process called *init* that is created during system initialization.
- ★ When a process terminates, the kernel changes the appropriate process descriptor pointers of all the existing children of the terminated process to make them become children of *init*.
- ★ This process monitors the execution of all its children and routinely issues `wait( )` system calls, whose side effect is to get rid of all zombies.

Process groups and login sessions

- ★ Modern Unix operating systems introduce the notion of *process groups* to represent a "job" abstraction.
- ★ For example, in order to execute the command line:

\$ ls | sort | more

a shell that supports process groups, such as bash, creates a new group for the three processes corresponding to ls, sort, and more.

- ★ In this way, the shell acts on the three processes as if they were a single entity i.e., job
- ★ Each process descriptor includes a *process group ID* field.
- ★ Each group of processes may have a *group leader*, which is the process whose PID coincides with the process group ID.
- ★ A newly created process is initially inserted into the process group of its parent.

Process groups and login sessions

- ★ Modern Unix kernels also introduce *login sessions*.
- ★ Informally, a login session contains all processes that are descendants of the process that has started a working session on a specific terminal—usually, the first command shell process created for the user.
- ★ All processes in a process group must be in the same login session.
- ★ A login session may have several process groups active simultaneously; one of these process groups is always in the foreground, which means that it has access to the terminal.
- ★ The other active process groups are in the background.
- ★ When a background process tries to access the terminal, it receives a SIGTTIN or SIGTTOU signal.
- ★ In many command shells the internal commands `bg` and `fg` can be used to put a process group in either the background or the foreground.

SOC 3010

OPERATING SYSTEM


WEEK 3 LECTURE 1

MEMORY MANAGEMENT

Memory Management

- ✱ Memory management is by far the most complex activity in a Unix kernel.
- ❑ **Virtual memory**
- ✱ All recent Unix systems provide a useful abstraction called *virtual memory*.
- ✱ Virtual memory acts as a logical layer between the application memory requests and the hardware Memory Management Unit (MMU).
- ✱ **Virtual memory has many purposes and advantages:**
 - Several processes can be executed concurrently.
 - It is possible to run applications whose memory needs are larger than the available physical memory.
 - Processes can execute a program whose code is only partially loaded in memory.
 - Each process is allowed to access a subset of the available physical memory.
 - Processes can share a single memory image of a library or program.
 - Programs can be relocatable, that is, they can be placed anywhere in physical memory.
 - Programmers can write machine-independent code, since they do not need to be concerned about physical memory organization.

Memory Management



- * The main ingredient of a virtual memory subsystem is the notion of *virtual address space*.
- * The set of memory references that a process can use is different from physical memory addresses.
- * When a process uses a virtual address, the kernel and the MMU cooperate to locate the actual physical location of the requested memory item.
- * These addresses have different nomenclatures depending on the computer architecture. Intel 80x86 manuals refer to them as "logical addresses".
- * Today's CPUs include hardware circuits that automatically translate the virtual addresses into physical ones.
- * To that end, the available RAM is partitioned into *page frames* 4 or 8 KB in length, and a set of page tables is introduced to specify the correspondence between virtual and physical addresses.
- * These circuits make memory allocation simpler, since a request for a block of contiguous virtual addresses can be satisfied by allocating a group of page frames having noncontiguous physical addresses.



Random access memory usage

- ★ All Unix operating systems clearly distinguish two portions of the random access memory (RAM).
- ★ A few megabytes are dedicated to storing the kernel image (i.e., the kernel code and the kernel static data structures).
- ★ The remaining portion of RAM is usually handled by the virtual memory system and is used in three possible ways:
 - To satisfy kernel requests for buffers, descriptors, and other dynamic kernel data structures
 - To satisfy process requests for generic memory areas and for memory mapping of files
 - To get better performance from disks and other buffered devices by means of caches

Random access memory usage

- ★ Each request type is valuable.
- ★ On the other hand, since the available RAM is limited, some balancing among request types must be done, particularly when little available memory is left.
- ★ Moreover, when some critical threshold of available memory is reached and a page-frame reclaiming algorithm is invoked to free additional memory, which are the page frames most suitable for reclaiming?
- ★ The only available solution lies in developing carefully tuned empirical algorithms.
- ★ One major problem that must be solved by the virtual memory system is *memory fragmentation*.
- ★ Ideally, a memory request should fail only when the number of free page frames is too small.
- ★ However, the kernel is often forced to use physically contiguous memory areas, hence the memory request could fail even if there is enough memory available but it is not available as one contiguous chunk.

Kernel Memory Allocator

- ★ The Kernel Memory Allocator (KMA) is a subsystem that tries to satisfy the requests for memory areas from all parts of the system.
- ★ Some of these requests will come from other kernel subsystems needing memory for kernel use, and some requests will come via system calls from user programs to increase their processes' address spaces.
- ★ A good KMA should have the following features:
 - It must be fast. Actually, this is the most crucial attribute, since it is invoked by all kernel subsystems (including the interrupt handlers).
 - It should minimize the amount of wasted memory.
 - It should try to reduce the memory fragmentation problem.
 - It should be able to cooperate with the other memory management subsystems in order to borrow and release page frames from them.

Kernel Memory Allocator

Several kinds of KMAs have been proposed, which are based on a variety of different algorithmic techniques, including:

- Resource map allocator
 - Power-of-two free lists
 - McKusick-Karels allocator
 - **Buddy system**
 - Mach's Zone allocator
 - Dynix allocator
 - **Solaris's Slab allocator**
- ★ **Linux's KMA uses a Slab allocator on top of a Buddy system.**

SOC 3010

OPERATING SYSTEM

WEEK 3 LECTURE 1

PROCESS VIRTUAL ADDRESS SPACE

Process virtual address space handling

- ★ The address space of a process contains all the virtual memory addresses that the process is allowed to reference.
- ★ The kernel usually stores a process virtual address space as a list of *memory area descriptors*.
- ★ For example, when a process starts the execution of some program via an `exec()`-like system call, the kernel assigns to the process a virtual address space that comprises memory areas for:
 - The executable code of the program
 - The initialized data of the program
 - The uninitialized data of the program
 - The initial program stack (that is, the User Mode stack)
 - The executable code and data of needed shared libraries
 - The heap (the memory dynamically requested by the program)


Process virtual address space handling

- * All recent Unix operating systems adopt a memory allocation strategy called *demand paging*.
- * With demand paging, a process can start program execution with none of its pages in physical memory.
- * As it accesses a nonpresent page, the MMU generates an exception; the exception handler finds the affected memory region, allocates a free page, and initializes it with the appropriate data.
- * In a similar fashion, when the process dynamically requires some memory by using `malloc()` or the `brk( )` system call (which is invoked internally by `malloc()`), the kernel just updates the size of the heap memory region of the process.
- * A page frame is assigned to the process only when it generates an exception by trying to refer its virtual memory addresses.
- * Virtual address spaces also allow other efficient strategies, such as the Copy-On-Write strategy.
- * For example, when a new process is created, the kernel just assigns the parent's page frames to the child address space, but it marks them read only.
- * An exception is raised as soon the parent or the child tries to modify the contents of a page.
- * The exception handler assigns a new page frame to the affected process and initializes it with the contents of the original page.

Swapping and caching

- ★ In order to extend the size of the virtual address space usable by the processes, the Unix operating system makes use of *swap areas* on disk.
- ★ The virtual memory system regards the contents of a page frame as the basic unit for swapping.
- ★ Whenever some process refers to a swapped-out page, the MMU raises an exception.
- ★ The exception handler then allocates a new page frame and initializes the page frame with its old contents saved on disk.

Swapping and caching



- ★ On the other hand, physical memory is also used as cache for hard disks and other block devices.
- ★ This is because hard drives are very slow: a disk access requires several milliseconds, which is a very long time compared with the RAM access time.
- ★ Therefore, disks are often the bottleneck in system performance.
- ★ As a general rule, one of the policies already implemented in the earliest Unix system is to defer writing to disk as long as possible by loading into RAM a set of disk buffers corresponding to blocks read from disk.
- ★ The sync() system call forces disk synchronization by writing all of the "dirty" buffers (i.e., all the buffers whose contents differ from that of the corresponding disk blocks) into disk.
- ★ In order to avoid data loss, all operating systems take care to periodically write dirty buffers back to disk.




SOC 3010

OPERATING SYSTEM

WEEK 3 LECTURE 1

DEVICE DRIVERS

Device Drivers




- * The kernel interacts with I/O devices by means of *device drivers*.
- * Device drivers are included in the kernel and consist of data structures and functions that control one or more devices, such as hard disks, keyboards, mice, monitors, network interfaces, and devices connected to a SCSI bus.
- * Each driver interacts with the remaining part of the kernel (even with other drivers) through a specific interface.
- * This approach has the following advantages:
 - Device-specific code can be encapsulated in a specific module.
 - Vendors can add new devices without knowing the kernel source code: only the interface specifications must be known.
 - The kernel deals with all devices in a uniform way and accesses them through the same interface.
 - It is possible to write a device driver as a module that can be dynamically loaded in the kernel without requiring the system to be rebooted. It is also possible to dynamically unload a module that is no longer needed, thus minimizing the size of the kernel image stored in RAM.



Device Drivers

* Figure showing how device drivers interface with the rest of the kernel and with the processes.



- * Some user programs (P) wish to operate on hardware devices.
- * They make requests to the kernel using the usual file-related system calls and the device files normally found in the */dev* directory.
- * Actually, the device files are the user-visible portion of the device driver interface.
- * Each device file refers to a specific device driver, which is invoked by the kernel in order to perform the requested operation on the hardware component.

Device Drivers

- ★ It is worth mentioning that at the time Unix was introduced graphical terminals were uncommon and expensive, and thus only alphanumeric terminals were handled directly by Unix kernels.
- ★ When graphical terminals became widespread, ad hoc applications such as the X Window System were introduced that ran as standard processes and accessed the I/O ports of the graphics interface and the RAM video area directly.
- ★ Some recent Unix kernels, such as Linux, include limited support for some frame buffer devices, thus allowing a program to access the local memory inside a video card through a device file.

SOC 3010 OPERATING SYSTEMS

Reading Assignments

Chapter 1 of Text Book

Text Book

➤ Understanding the Linux Kernel

- Daniel P/ Bovet and Marco Cesati
Edition, O' Reilly Media 2005, ISBN 978-0596005658**

SOC 3010

OPERATING SYSTEM

END OF
WEEK 3 LECTURE 1

UNIX/LINUX KERNELS

WISH YOU ALL THE BEST

INHA UNIVERSITY TASHKENT

FALL SEMESTER 2024

SOC 3010

OPERATING SYSTEM

CREDITS/HOURS PER WEEK : 3/3

COURSE TYPE : TECHNICAL CORE SEQUENCE

INSTRUCTOR

DR. A. R. NASEER

HEAD & PROFESSOR OF COMPUTER SCIENCE & ENGG

